

**Database Systems — CSci 4380**  
**Midterm Exam #2**  
**November 1, 2018**

**SOLUTIONS**

**Question 1.** Write the following queries using **SQL** for the data model below. The data model is described in detail in the back of the exam.

Similar to Homeworks #4 and #5, your solutions should be a single SQL expression. However, you can use any SQL construct that you wish. Make sure that your queries are easily readable and syntactically correct.

- (a) (10 points) Return the id of all trips that either started or ended in 2017 or the trip had an associated message that was sent in 2017. (Hint: remember `extract (year from datevalue)`).

**Solutions:**

---

```
select
    id
from
    trips t1
where
    extract(year from t1.fromdate) = 2017
    or extract(year from t1.todate) = 2017
union
select
    tripid
from
    messages
where
    extract(year from sentdatetime) = 2017;
```

---

- (b) (10 points) Find and return the name, state, city of all users who stayed in the same house as customer in two different trips with start dates at least 90 days apart and reviewed the earlier trip with `isclean = 1`.

**Solutions:**

---

```
select distinct
    u.name
    , u.state
    , u.city
from
    users u
    , trips t1
    , trips t2
    , reviews r
where
    u.email = t1.customeremail
    and u.email = t2.customeremail
    and t1.houseid = t2.houseid
    and t2.fromdate - t1.fromdate >= 90
    and r.tripid = t1.tripid
    and r.isclean = 1

-- having the following is not wrong, but not required for full points
    and t1.isconfirmed = True and t2.isconfirmed=True
;
```

---

- (c) (12 points) Return the name of amenities that are listed for at least three houses in New York state such that there is a confirmed trip in the database for each house. These amenities should not appear in a house listed in any other state.

**Solutions:**

---

```
select
    ha.amenity
from
    houseamenities ha
    , houses h
    , trips t
where
    ha.houseid = h.id
    and state = 'New York'    --- also fine to use 'NY'
    and t.houseid = h.id
    and t.inconfirmed = True
group by
    ha.amenity
having
    count(distinct h.id)>=3
except
select
    distinct ha.amenity
from
    houseamenities ha
    , houses h
where
    ha.houseid = h.id
    and h.state <> 'New York';

alternate
```

```
select
    ha.amenity
from
    houseamenities ha
    , houses h
    , trips t
where
    ha.houseid = h.id
    and state = 'New York'    --- also fine to use 'NY'
    and t.houseid = h.id
    and t.inconfirmed = True
    and not exists (
        select 1
        from
            houseamenities ha1
            , houses h1
        where
            ha1.houseid = h1.id
            and h1.state <> 'New York'
            and ha1.amenity = ha.amenity)
group by
    ha.amenity
having
    count(distinct h.id)>=3 ;
```

or

```
select
```

```

    ha.amenity
from
    houseamenities ha
    , houses h
    , trips t
where
    ha.houseid = h.id
    and state = 'New York'    --- also fine to use 'NY'
    and t.houseid = h.id
    and t.inconfirmed = True
    and ha.amenity not in (
        select ha1.amenity
        from
            houseamenities ha1
            , houses h1
        where
            ha1.houseid = h1.id
            and h1.state <> 'New York')
group by
    ha.amenity
having
    count(distinct h.id)>=3 ;

```

---

- (d) (12 points) Find all unconfirmed trips for a house in Troy, New York that were supposed to start on '10/31/2018'. For each trip, return the house id, house label, email of the owner as well as number of messages sent by the owner for this trip.

**Solutions:**

---

```

select
    h.id
    , h.label
    , h.owneremail
    , count(m.id) as numbyowner
from
    trips t
    join houses h on h.id = t.houseid
    left join messages m
        on t.id = m.tripid
        and m.senderemail = h.owneremail
where
    t.isconfirmed = False    -- this can also go in the join statement
    and h.state = 'New York'
    and h.city = 'Troy'
group by
    t.id    -- it is ok not to group by this as well
            -- if we assume one unconfirmed trip per day
    , h.id
    , h.label, h.owneremail ; -- technically we do not need to join by these
                                -- two but adding them is not wrong either

```

note: we will also accept

```
and m.senderemail = t.customeremail
```

instead of

```
and m.senderemail = h.owneremail
```

---

- (e) (10 points) Delete all reviews entered by users who are also owners of houses. Such users should have no business reviewing others.

## Solutions:

---

```
delete from reviews where tripid in (select t.tripid from trips t,  
houses h where t.customeremail = h.owneremail) ;
```

---

**Question 2 (14 points).** Write a trigger that activates after a tuple is inserted into the `Reviews` table. Recall that each review is for a trip and each trip is for a specific house.

The trigger updates the tuple in relation `Houses` for this house and sets the average rating attribute (`Houses.avgrating`) to the average of the `Reviews.isgoodvalue` values for all trips involving this house.

Complete the trigger description given below for simplicity. You are free to use any pl/pgsql type procedural code here.

**Solutions:**

---

```
CREATE FUNCTION avgreview_f () RETURNS trigger AS $$
DECLARE
    val FLOAT ;
    idval int ;
BEGIN

SELECT houseid INTO idval FROM trips WHERE id = NEW.id;

SELECT avg(isgoodvalue) into val
FROM reviews r, trips t
WHERE
    r.tripid = t.id
    and t.houseid = idval ;

UPDATE houses
SET avgval = val
WHERE houseid = idval ;

RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER avgreview_trigger AFTER INSERT ON reviews
FOR EACH ROW EXECUTE PROCEDURE avgreview_f();
```

---

Write your answers in the box below only. Do not write on the back or outside the box.

---

**Question 3 (a-10 points).** Create a view called `housestats` that returns for each house in the `Houses` table the following attributes: `Houses.id`, `Houses.label`, `Houses.state`, `Houses.city`, `numtrips` (total number of confirmed trips), `numdaysrented`=(total number of days the house is rented by the confirmed trips) by completing the following expression:

**Solutions:**

---

```
create view housestats(id, label, state, city, numtrips,
numdaysrented) as
select
    h.id
    , h.label
    , count(*) as numtrips
    , sum(todate-fromdate) as totalduration
from
    trips t
    , houses h
where
    t.houseid = h.id
    and t.isconfirmed = True
group by
    h.id ;
```

---

**Question 3 (b-10 points).** Use your view from Question 3(a) above to find the most popular houses in Troy, New York based on `numtrips`. For these houses, return all the attributes in the view `housestats`.

**Solutions:**

---

```
select
    *
from houseinfo h
where
    h.state = 'NY'
    and h.city = 'Troy'
    and h.numtrips = (
        select max(h2.numtrips)
        from houseinfo h2
        where h2.state = 'NY'
        and h2.city = 'Troy');
```

---

**Question 4 (12 points).** You are given the following table definitions and table contents.

```
CREATE TABLE ROLES(id INT PRIMARY KEY, name VARCHAR(10) NOT NULL, type VARCHAR(10)) ;
CREATE TABLE SHOWS(id INT PRIMARY KEY, title VARCHAR(100) NOT NULL);
CREATE TABLE APPEARS(id INT PRIMARY KEY, sid INT, rid INT NOT NULL
    , FOREIGN KEY (sid) REFERENCES SHOWS(id) ON DELETE CASCADE ON UPDATE SET NULL
    , FOREIGN KEY (rid) REFERENCES ROLES(id) ON DELETE CASCADE ON UPDATE SET NULL);
```

Roles			Shows		Appears		
id	name	type	id	title	id	sid	rid
1	Rick	Rick	5	Meeseeks and Destroy	1	5	1
2	Morty	Morty	12	A Rickle in Time	2	5	2
3	PickleRick	Rick	24	Pickle Rick	3	24	1
4	Evil Rick	Rick	28	The Ricklantis Mixup	4	24	2
5	Rick Morty	Morty			5	24	3
6	Jaguar				6	24	6
					7	28	1
					8	28	2
					9	28	5

For each operation below, describe which rows from which tables are changed/deleted and why (or why not). Assume each operation operates on the table contents listed above (hence each part is independent).

- DELETE FROM roles WHERE name = 'Jaguar' ;
- UPDATE shows SET id = 24 WHERE title LIKE '%Rick%';
- UPDATE roles SET id = 8 WHERE id = 5;
- DELETE FROM shows WHERE NOT EXISTS (SELECT \* FROM appears WHERE sid=shows.id);

**Solutions:**

- Delete (6,Jaguar) from roles, and (6,24,6) from Appears because of cascade.
- Try to change the id for both 24 and 28 to 24, which cannot be due to primary key constraint. The update fails and no tables are changed.
- Try to change id 5 to id 8, which results in a SET NULL for rid of tuple (9,28,5). But, rid cannot be null, so the whole transaction fails. This results in no changes to the tables.
- Delete (12, A Rickle in Time) from shows. No changes to other tables.

Use this page for scratch work only. Do not share your solutions or any drafts of your solutions with anyone.



## Data model to be used in Exam #2

This is a data model based on the E-R problem from Homework#3, storing information about house rentals in a system similar to Airbnb. It is simplified in various ways to make it easier to use in an exam. Please read carefully. The keys of each relation are underlined.

### **Users(email, password, name, street, state, city, zip, country, ccno)**

Stores information about all the users in the system, renters and customers including credit card number (ccno).

### **Houses(id, label, description, street, state, city, zip, country, price, owneremail, avgrating)**

Each house is listed by an owner and we store the email of the owner for each house. The daily price of a house is fixed for simplicity. Average rating is the average rating value for this house.

### **HouseAmenities(houseid, amenity)**

This relation lists the various amenities a house has such as 'free wifi', 'garage parking', etc,

### **Trips(id, fromdate, todate, isconfirmed, totalprice, customeremail, houseid)**

Each trip has a start and end date, whether it is confirmed or not (true/false value) and a total price for the trip. Each trip is by a user (customer), so his/her email is stored. This is the person paying for the trip. Each trip is for a specific house, so the id of the house is also stored. (The renter can be deduced from the owner of the house).

### **Reviews(tripid, isclean, isgoodvalue, isgoodrenter, reviewtext)**

We will only store the reviews of the houses in this database. There can only be one review for each trip, so we will use tripid as the key. Each review is for a specific house and its renter (stored in other relations) and contains 1-5 star values for isclean, isgoodvalue, isgoodrenter as well as a review text.

### **Messages(id, senderemail, receiveremail, messagetext, sentdatetime, prevmid, tripid)**

Finally, we store messages from one user (sender) to another user (receiver) together with the text and datetime it was sent. The attribute **prevmid** can be null if this is an original email or store the id of the message that this message was sent in response to. Each message is for a specific trip.