

**Exam 2 Answers, Spring 2013**  
**CSCI 4380 Database Systems**  
**Time: 110 minutes**

**Name :** \_\_\_\_\_

<b>Q1.</b>	<b>Q4.</b>
<b>Q2.</b>	
<b>Q3.</b>	<b>Total.</b>

**Rules.** Open book and notes. Do not use any electronic tools including your computer. Work alone. You **cannot** talk to anyone in class, or share notes or thoughts.

**Question 1.** Answer all parts of this question using the data model in the appendix. Write the following queries using SQL only (no procedural code is allowed). Make sure you use DISTINCT only if you have to.

- (a) **(14 points)** Find all people who are friends with 'Clara Oswin Oswald' both on Facebook and Twitter since 2012. Return their names.

```
SELECT DISTINCT
  p1.name
FROM
  people p1
  , tw tf
  , fb ff
  , people p2
WHERE
  p1.name = 'Clara Oswin Oswald'
  and p1.tuser = tf.user
  and tf.whenlinked >= date '01-01-2012'
  and ( (p1.fuser = ff.user1 and ff.user2 = tf.friend)
        or (p1.fuser = ff.user2 and ff.user1 = tf.friend)
      )
  and tf.friend = p2.tuser
  and tf.whenlinked >= date '01-01-2012';
```

- (b) **(14 points)** Find all people for whom the number of followers on Twitter is at least 10 times the number of friends on Twitter. Return the name of the people, the number of followers and the number friends on Twitter.

If the person has no friends, then return him if he has at least 10 followers. (**Hint: you can do this query with a single SELECT/FROM/WHERE/GROUP BY/HAVING clause with no nested subqueries.**)

```
SELECT
    p.name
    , count(distinct tfo.user) as numFollower
    , count(distinct tfr.friend) as numFriend
FROM
    people p
    (join tw tfo
        on p.tuser = tfo.friend)
    left outer join tw tfr
        on p.tuser = tfr.user
GROUP BY
    p.fid
    , p.name
HAVING
    count(distinct tfo.user) >= 10* count(distinct tfr.friend)
or ( count(distinct tfo.user) >= 10
    and count(distinct tfr.friend) = 0 )
```

- (c) **(14 points)** Find pairs of people who are friends on Facebook, but are not friends on Twitter. Return their ids. (You can return one or two tuples for each pair of ids).

```
select
    p1.fid
    , p2.fid
from
    people p1
    , fb ff
    , people p2
where
    (p1.fuser = ff.user1
    and p2.fuser = ff.user2)
    or
    (p2.fuser = ff.user1
    and p1.fuser = ff.user2)
    and not exists
        (select
            1
        from
            tw t1
        where
            p1.tuser = t1.user
            and p2.tuser = t2.follower
        ) ;
```

- (d) **(10 points)** Write at most two update statements in a transaction block (no other procedural code is allowed) to update the type attribute for people to 'active' if they have exchanged a message on Twitter in the last month, and 'passive' otherwise.

```
begin transaction ;
update people set type = 'passive' ;
update people set type = 'active'
where tuser in (select t.user from tw
                where now() - lastcommunicated <= interval '1 month');
end ;
```

**Question 2 (18 points).** We are now mining Facebook for some information. Suppose you are given the new table below:

```
CREATE TABLE fb_stats (  
    state          VARCHAR(100)  
    , agegroup      INT PRIMARY KEY  
    , avg_ff_num    FLOAT  
    , earliest_adoption DATE  
    , PRIMARY KEY (state, agegroup)  
) ;
```

This table stores the earliest adoption date of Facebook of all people in that state and age group. The earliest adoption date for a person is the earliest day that they were linked to someone on Facebook (given by the `whenlinked` attribute).

The table also stores the average number of friends per person (`avg_ff_num`) in Facebook of all people in the same state and age group.

Write code in pseudo-code to populate this table based on data available in the database. You can use a single SQL query or any procedural code, however you must spell out any query in proper SQL.

Age group is a value of the form 10, 20, 30, etc. to denote people with ages 0-10, ages 11-20, ages 21-30, etc. Here is some help. The code below finds the age of a person born on January 1, 1985 (by finding the number of days between now and his birth day):

```
select trunc(extract(days from now() - date '01-01-1985')/365) ;
```

```
INSERT INTO fb_stats  
SELECT  
    tmp.state,  
    , 10 * trunc(extract(days from now() - tmp.dateofbirth)/3650 as int)  
    , avg(numfr)  
    , min(earliest)  
FROM (  
    SELECT  
        p.fid  
        , p.dateofbirth  
        , min(p.whenlinked) as earliest  
        , count(*) as numfr  
    FROM  
        people p  
        , fb fr  
    WHERE  
        p.fuser = fr.user1  
        or p.fuser = fr.user2  
    GROUP BY  
        p.fid  
        , p.dateofbirth  
        , p.state  
    ) as tmp  
GROUP BY  
    tmp.state,  
    10 * trunc(extract(days from now() - tmp.dateofbirth)/3650 as int) ;
```

**Question 3 (10 points).** Assume you are already given the table `fb_stats` populated with the statistics from Question 2. Write the following query using SQL with or without procedural components.

We are computing some measure of hipness. For each agegroup, find the state with the earliest adoption date. Return the agegroup, the state with the earliest adoption date and the adoption date. If multiple states are tied for a specific age group, return all of them.

```
SELECT
    f.agegroup
    , f.state
FROM
    fb_stats f
WHERE
    f.earliest_adoption <= ALL
        ( SELECT
            f2.earliest_adoption
          FROM
            fb_stats f2
          WHERE
            f2.agegroup = f.agegroup )
ORDER BY
    f.agegroup
    , f.state
```

**Question 4 (20 points).** You are given the following table definitions and instances. Each of the statements below operate on the original tables. Write down the results of each operation below by listing only the changed tuples. Provide a short sentence of why these tuples were changed.

```
CREATE TABLE d (
    id      INT PRIMARY KEY
    , name VARCHAR(100) );
CREATE TABLE e (
    id      INT PRIMARY KEY
    , name VARCHAR(100)
    , did   INT NOT NULL FOREIGN KEY
          REFERENCES d(id) ON UPDATE CASCADE
    , pcnt  INT ) ;
CREATE TABLE ep (
    pid     INT
    , eid    INT
    , PRIMARY KEY(pid, eid)
    , FOREIGN KEY (pid) REFERENCES p(id)
      ON DELETE CASCADE ON UPDATE CASCADE
    , FOREIGN KEY (eid) REFERENCES e(id)
      ON DELETE CASCADE ON UPDATE CASCADE ) ;
```

```
CREATE TABLE p (
    id      INT PRIMARY KEY
    , name VARCHAR(100) ) ;
CREATE TRIGGER epdel AFTER DELETE ON ep
FOR EACH ROW
REFERENCING OLD ROW AS old
DECLARE
    c int ;
BEGIN
    SELECT count(*) INTO c
    FROM ep
    WHERE pid = old.pid ;
    UPDATE e SET pcnt = (SELECT count(*)
                          FROM ep WHERE ep.eid = old.eid);
    IF c <= 1 THEN
        DELETE FROM p WHERE id = old.pid ;
    END IF ;
END ;
```

id	name	id	name	did	pcnt	id	name	pid	eid
1	'da'	1	'ea'	1	1	1	'pa'	1	1
2	'db'	2	'eb'	1	2	2	'pb'	1	2
3	'dc'	3	'ec'	2	0	3	'pc'	1	4
		4	'ed'	1	2			2	2
								2	4
(d)		(e)		(p)		(ep)			

(a) DELETE FROM d WHERE d.name = 'db';

**Answer.** Fails because there is no rule on e for delete.

(b) DELETE FROM ep WHERE ep.pid = 2 and ep.eid=4;

**Answer.** ep: (1,1), (1,2), (1,4), p: project 2 deleted, e: all pcnt for 2 and 4 are set to 1.

(c) UPDATE e SET id = 22 WHERE id = 2;

**Answer.** ep changed to (1,1), (1,22), (1,4), (2,22), (2,4), and the id for 2 is changed to 22 in e.

(d) DELETE FROM e WHERE id < 3;

**Answer.** ep empty, p only 3 left. e: only 3/4 left, the data remains the same.

**Question 5 (0 points).** What is good movie pitch based on SQL? Here is my pitch for “Expandables”:

```
select
    distinct actors
from
    blockbusters
where
    year >= 1990
    and year < 2000
    and nr_order = 1
limit
    8
```

**Appendix.** This is a data model for keeping track of a set of people of interest to you in two different social networks, Facebook and Twitter.

```
CREATE TABLE people (  
    fid          INT PRIMARY KEY  
    , fuser      VARCHAR(20)  UNIQUE  
    , tuser      VARCHAR(20)  UNIQUE  
    , name       VARCHAR(100) NOT NULL  
    , dateofbirth DATE  
    , state      VARCHAR(100)  
    , type       VARCHAR(20)  
) ;
```

Stores all the people in the database. Attributes **fuser** and **tuser** refer to the username of the user on Facebook and Twitter respectively. The **state** attribute stores where the person lives (if known), assuming of course that all the people in this database are from US. The **type** attribute will be described in one of the questions.

```
CREATE TABLE fb (  
    user1        VARCHAR(20)  
    , user2      VARCHAR(20)  
    , how        VARCHAR(100)  
    , whenlinked DATE  
    , lastcommunicated DATE  
    , PRIMARY KEY(user1,user2)  
    , FOREIGN KEY (user1) REFERENCES person(fuser)  
      ON DELETE CASCADE ON UPDATE CASCADE  
    , FOREIGN KEY (user2) REFERENCES person(fuser)  
      ON DELETE CASCADE ON UPDATE CASCADE  
) ;
```

```
CREATE TABLE tw (  
    user         VARCHAR(20)  
    , friend     VARCHAR(20)  
    , whenlinked DATE  
    , lastcommunicated DATE  
    , PRIMARY KEY(user, friend)  
    , FOREIGN KEY (user) REFERENCES person(tuser)  
      ON DELETE CASCADE ON UPDATE CASCADE  
    , FOREIGN KEY (friend) REFERENCES person(tuser)  
      ON DELETE CASCADE ON UPDATE CASCADE  
) ;
```

The table **fb** stores who is friends with whom in Facebook, and the table **tw** stores the same in Twitter. Friendship in Facebook is symmetric, but only one tuple per pair is stored. If (**user1**,**user2**) is in **fb**, then **user1** friended **user2**, and **user2** approved it.

For Twitter, relationships are not necessarily symmetric. If A (**user**) is a friend of B (**friend**) on Twitter, then it means that B is a follower of A on Twitter. If A and B are mutual friends , then two tuples are stored.

For both tables, we store the date when the pair last communicated: i.e. they exchanged one message in each direction (A to B, and B to A) on the same day. Note that for two people to communicate on Twitter, they must be mutual friends.