

Base Systems — CSci 4380  
 Midterm Exam #2  
 March 25, 2018

# SOLUTIONS

**Question 1.** Write the following queries using SQL using the data model below. The model is described in detail in the back of the exam.

In all queries, use **DISTINCT** only if you have to. Do not use **ORDER BY** unless a specific ordering is asked. Write your queries in a readable format.

```
Users(username, name, email, password, address)
FriendsWith(username1, username2, sincewhen)
Landmarks(id, landmarkname, landmarktype, state, city, zip, country)
Segments(id, startx, starty, endx, endy, landmark_id)
Events(id, username, eventtype, eventdate, starttime, endtime)
DataPoints(id, event_id, seqno, starttime, endtime, segment_id)
Comments(id, username, event_id, comment_text, whenposted)
```

- (a) **(12 points)** Return the id, name of all landmarks in `Troy` (city) `NY` (state) that users pass by in `cycling` or `skating` events (eventtype).

**Solutions:**

```
select distinct
    l.id
    , l.name
from
    events e
    , datapoints d
    , segments s
    , landmarks l
where
    e.id = d.event_id
    and e.eventtype in ('cycling', 'skating') -- can also use an OR here
    and d.segment_id = s.id
    and s.landmark_id = l.id
    and l.city = 'Troy'
    and l.state='NY' -- 'New York' is also fine
```

- (b) **(10 points)** For each user, return the number of events they participated in and the duration of their longest duration event (duration is `events.endtime-events.starttime`). Note that the number of events can be zero for a user, in which case you will return null for the duration value of this user.

**Solutions:**

---

```
select
  u.username
  , count(e.username)          -- must count an attribute
                              -- otherwise will not get zero here
  , max(e.endtime-e.starttime) -- will be null if no events for user
from users u left join events e -- left join is required
  on u.username= e.username
group by u.username
```

---

- (c) **(10 points)** Return the `username` and `email` of all users who have commented on 4 or more different events (`event_id`) that were created on the same date (`eventdate`).

**Solutions:**

---

```
select distinct          -- distinct is needed here, otherwise we
    u.username          -- will return a tuple for each matching date
    , u.email
from
    users u
    , events e
    , comments c
where
    u.username = c.username
    and e.id = c.event_id
group by
    u.username
    , u.email            -- this can be omitted, email is unique given username
    , e.eventdate        -- group by date to find number of events
                        -- with a comment on that day
having
    count(distinct e.id) >= 4; -- distinct is needed here because
                                -- there could be multiple comments
```

---

- (d) **(12 points)** We are going to sell some user data to a third party for political targeting. Return the email of all users who either have a **powerwalking** event (eventtype) themselves or are friends with a person who has a **powerwalking** event (eventtype) in the database.

**Solutions:**

---

```
select
    u.email
from
    users u
    , events e
where
    u.username = e.username
    and e.eventtype = 'powerwalking'
union
select
    u.email
from
    users u
where
    exists (select 1
            from
                friends f
                , events e
            where
                e.eventtype = 'powerwalking'
                and ((f.username1=u.username and f.username2=e.username)
                    or
                    (f.username2=u.username and f.username1=e.username))

select
    u.email          -- distinct is not needed because of union
from
    users u          -- but no performance/score penalty if it is there
    , events e       -- as union will do duplicate removal anyway
where
    u.username = e.username
    and e.eventtype = 'powerwalking'
union
select
    u.email
from
    users u
    , events e
    , friends f
where
    ((u.username = f.username1
     and f.username2 = e.username
     or
     (u.username = f.username2
     and f.username1 = e.username))
    and e.eventtype = 'powerwalking')
```

---

- (e) **(10 points)** Delete all landmarks from the database with no associated segment.

**Solutions:**

---

```
delete from
  landmarks l
where
  not exists
    (select 1    --does not matter what is returned in not exists
     from
       segments s
     where
       s.landmark_id = l.id);
```

---

- (f) **(14 points)** Return pairs of segment id and username for each segment in the database such that the username belongs to a user with the fastest running time (smallest **endtime-starttime**) for that segment. Remember, multiple users may have the same running time; we will return them all in different tuples.

**Solutions:**

---

```
select distinct
    d.segment_id
    , u.username
from
    users u
    , events e
    , datapoints d
where
    u.username = e.username
    e.eventtype = 'running' -- can also be omitted as question is ambiguous
    and e.id = d.event_id
    and not exists
        (select 1
         from
             datapoints d2
             , events e2
         where
             e2.id = d2.event_id
             and e2.eventtype = 'running'
             and e2.username = u.username
             and d2.segment_id = d.segment_id
             and (d2.endtime - d2.starttime) < (d.endtime - d.starttime);
```

---

**Question 2 (16 points).** For this question only, you can use a single expression, or you can piece together multiple queries, inserts and auxiliary tables for this question. You do not have to put them inside a procedure block and you do not need to drop your auxiliary tables.

We are running a promotion to send energy drinks with drones to a select set of users in the database just before their next event.

To facilitate this, find and return the `username` of highly predictable users in the database and the `segment_id` of their most popular starting location. The starting location is the `segment_id` for the `datapoint` with `seqno=1` for that event. Returned users must have had an event in at least 300 days of 2017. The returned `segment_id` for this user must be their starting location in 90% or more of all his/her events in the database.

### Solutions:

---

```
create table u as --users with events in 300+ days in 2017
```

```
select
from
    users u
    , events e
where
    u.username = e.username
    and e.eventdate > '1/1/2017'
    and e.eventdate <= '12/31/2017'
group by
    u.username
having
    count(distinct e.eventdate)>= 300 ;
```

```
select
    u.username
    , e.segment_id
from
    events e
    , u
    , datapoints d
where
    e.username = u.username
    and d.event_id=e.id
    and d.seqno=1 -- startnig segment
group by
    u.username
    , d.segment_id
having count(*) >=
    0.9* (select count(*)
        from events e2
        where e2.username=u.username)
```

---

**Question 3 (16 points).** You are given the following table definitions and instances. For each operation, show the changes to the tables by directly drawing on the tables. Provide a short sentence of why these tuples were changed or not changed right below the query.

```
CREATE TABLE abc (
    id INT PRIMARY KEY, name CHAR(2) );

CREATE TABLE def (
    id INT PRIMARY KEY, key INT NOT NULL) ;

CREATE TABLE ghi (
    id1 INT, id2 INT, val INT
    , PRIMARY KEY (id1, id2)
    , FOREIGN KEY (id1) REFERENCES abc(id)
      ON UPDATE CASCADE
    , FOREIGN KEY (id2) REFERENCES def(id)
      ON UPDATE CASCADE ) ;

CREATE TABLE jkl (
    id INT PRIMARY KEY, id2 INT, val INT
    , FOREIGN KEY (id2) REFERENCES def(id)
      ON DELETE CASCADE ON UPDATE CASCADE);
```

```
create function e2f(x int) returns void AS $$
DECLARE
    c INT ;
BEGIN
    BEGIN; -- start a transaction block
        SELECT count(*) INTO c FROM abc WHERE id=x;
        DELETE FROM abc WHERE id = x;
        IF c>0 THEN
            INSERT INTO def(id) VALUES(x);
        END IF;
    COMMIT;
END ; $$ LANGUAGE plpgsql ;

CREATE TRIGGER fixit BEFORE DELETE ON def
FOR EACH ROW REFERENCING NEW ROW AS NEW
BEGIN
    UPDATE ghi SET id2 =
        (SELECT min(id) FROM def WHERE key IS NOT NULL)
    WHERE id2 = NEW.id ;
END ;
```

(a) DELETE FROM abc WHERE name = 'joy';	<u>id</u>		<u>name</u>		<u>id</u>		<u>key</u>		<u>id1</u>	<u>id2</u>	<u>val</u>	<u>id</u>	<u>id2</u>	<u>val</u>
	1		joy		6		5		1	6	4	11	6	3
	2		nya		7		4		1	7	5	12	6	2
	3		sky		8		1		3	8	3	13	8	
	(abc)		(def)		(ghi)		(jkl)							
(b) DELETE FROM def WHERE key>4;	<u>id</u>		<u>name</u>		<u>id</u>		<u>key</u>		<u>id1</u>	<u>id2</u>	<u>val</u>	<u>id</u>	<u>id2</u>	<u>val</u>
	1		joy		6		5		1	6	4	11	6	3
	2		nya		7		4		1	7	5	12	6	2
	3		sky		8		1		3	8	3	13	8	
	(abc)		(def)		(ghi)		(jkl)							
(c) UPDATE def SET id=id*10 WHERE key=1;	<u>id</u>		<u>name</u>		<u>id</u>		<u>key</u>		<u>id1</u>	<u>id2</u>	<u>val</u>	<u>id</u>	<u>id2</u>	<u>val</u>
	1		joy		6		5		1	6	4	11	6	3
	2		nya		7		4		1	7	5	12	6	2
	3		sky		8		1		3	8	3	13	8	
	(abc)		(def)		(ghi)		(jkl)							
(d) SELECT e2f(3);	<u>id</u>		<u>name</u>		<u>id</u>		<u>key</u>		<u>id1</u>	<u>id2</u>	<u>val</u>	<u>id</u>	<u>id2</u>	<u>val</u>
	1		joy		6		5		1	6	4	11	6	3
	2		nya		7		4		1	7	5	12	6	2
	3		sky		8		1		3	8	3	13	8	
	(abc)		(def)		(ghi)		(jkl)							



**Solutions:**

- a.** There is no delete cascade on ghi for id1. As a result, the delete on abc fails. No tables are changed.
- b.** Deleting (6,5) from def, trigger is activated, in ghi (1,6,4) is changed to (1,6,4). There is no delete cascade on ghi and a referencing tuple. So delete fails.
- c.** Changing (8,1) to (80,1) in def. There is update cascade on both ghi and jkl. So, they are both changed ((3,8,3) to (3,80,3), (13,8,null) to (13,80,null)).
- d.** select count(\*) sets c to 1. We try to delete (3,'sky') from abc but it fails because there is no delete cascade on ghi for id1.

## Data model to be used in Exam #2

This is a slightly modified version of the data model from Exam #1. The main change is the primary key for landmarks now is an id, allowing multiple landmarks of the same name. We have also added comments.

```
create table users ( -- all users in the system
    username      varchar(12) primary key
    , name        varchar(100) not null
    , email       varchar(100) not null
    , password    varchar(100) not null
    , address     varchar(100)
) ;

create table friendswith (
    -- friendship is mutual, but stored in one direction only,
    -- username1 is the person initiated the friendship
    username1     varchar(12)
    , username2   varchar(12)
    , sincewhen   date -- when friendship was confirmed
    , primary key (username1, username2)
    , foreign key (username1) references users(username)
    , foreign key (username2) references users(username)
) ;

create table landmarks (
    id            int primary key
    , landmarkname varchar(100) not null
    , landmarktype varchar(100) --e.g. building, monument, etc.
    , state       varchar(100)
    , city        varchar(100)
    , zip         varchar(20)
    , country     varchar(100)
) ;

create table segments (
    id            int primary key
    , startx      numeric(8,4) not null
    , starty      numeric(8,4) not null
    , endx        numeric(8,4) not null
    , endy        numeric(8,4) not null
    , landmark_id int
    , foreign key (landmark_id) references landmarks(id)
) ;

create table events (
    id            int primary key
    , username     varchar(12) not null
    , eventtype    varchar(100) not null --cycling, running, etc.
    , eventdate    date not null
    , starttime    time not null
    , endtime      time
    , foreign key (username) references users(username)
) ;
```

```

create table datapoints (
    id                int primary key
    , event_id        int not null
    , seqno            int not null
    , starttime        time
    , endtime          time
    , segment_id       int
    , foreign key (event_id) references events(id)
    , foreign key (segment_id) references segments(id)
) ;

-- a data point is a segment on someone's running or cycling event,
-- the first data point has seqno=1, followed by seqno: 2,3,4 ...
-- describing the segments that one has passed during the event.
-- The end point of a segment for a data point with seqno x is the starting
-- point of the segment for the data point with seqno x+1.

create table comments (
    -- each comment is made by the user with given username
    -- for a specific event
    id                int primary key
    , username         varchar(12) not null
    , event_id         int not null
    , comment_text     text not null
    , whenposted       date
    , foreign key (username) references users(username)
    , foreign key (event_id) references events(id)
) ;

```