



Database Systems — CSci 4380

Final Exam

May 7, 2018

SOLUTIONS

Question 1. We are given the data model below from Exams 1 and 2. The company for this database wants to improve user engagement and came up with some questionable functionality built on the queries below.

Write the following queries using SQL. In all queries, use `DISTINCT` only if you have to. Do not use `ORDER BY` unless a specific ordering is asked. Write your queries in a readable format.

```
Users(username, name, email, password, address)
FriendsWith(username1, username2, sincewhen)
Landmarks(id, landmarkname, landmarktype, state, city, zip, country)
Segments(id, startx, starty, endx, endy, landmark_id)
Events(id, username, eventtype, eventdate, starttime, endtime)
DataPoints(id, event_id, seqno, starttime, endtime, segment_id)
Comments(id, username, event_id, comment.text, whenposted)
```

- (a) (8 points) For each user, return the `username` and `segment_id` of all segments involving at least three different `eventtypes` for events that took place in 2018.

Solutions:

```
select
    e.username
    , d.segment_id
from
    events e
    , datapoints d
where e.id and d.event_id
    and extract(year from e.eventdate) = 2018
group by
    e.username
    , d.segment_id
having
    count(distinct e.eventtype) >= 3;
```

- (b) (12 points) For the user with `username = 'TheFlash'`, return the `username` of all users who had an overlapping event with him in May 2018. In other words, according to `datapoints`, both users passed the same segment on the same day for their events with overlapping start and end times (you can simply check that the start time of one of the tuples was between the start and end times of the other tuple).

Solutions:

```
select distinct e2.username
from events e, datapoints d, datapoints d2, events e2
where e.username = 'TheFlash'
    and d.event_id = e.id
    and e.eventdate >= '1/5/2018'
    and e2.eventdate = e.eventdate
    and e2.username <> e.username
    and d2.event_id = e2.id
    and d2.segment_id = e2.segment_id
    and (d2.starttime <= d.endtime and d.starttime <= d2.starttime
        or d.starttime <= d2.endtime and d2.starttime <= d.starttime);
```

- (c) (14 points) For the user with `username = 'TheFlash'`, return the id of the top 5 most popular segments visited by his friends. Popularity is determined by the total number of data points for that segment by any friend of `TheFlash`. The returned segments should have never been visited by `TheFlash` (no data points).

Solutions:

```
select d.segment_id, count(*) as total
from friendswith f, events e, datapoints d
where ((f.username1 = 'Flash' and e.username = f.username2)
      or (f.username2 = 'Flash' and e.username = f.username1))
      and e.id = d.event_id
      and d.segment_id not in
          (select d2.segment_id from events e2, datapoints d2
           where e2.username = 'Flash' and d2.event_id = e2.event_id)
group by d.segment_id
order by total desc
limit 3;
```

Question 2. You are given the following schedule of operations:

$$r_3(Y) w_3(Y) r_1(X) r_1(Y) w_1(X) r_2(X) r_2(Z) w_3(Z) w_2(W)$$

- (a) **(10 points)** List all conflicts and draw the conflict graph.

Is this schedule serializable? If it is, show a serial order that is equivalent to this schedule. If it is not serializable, discuss why not.

Solutions: Conflicts: $w_3(Y) r_1(Y)$, $w_1(X) r_2(X)$, $r_2(Z) w_3(Z)$.

Graph, $T3 -> T1$, $T1 -> T2$, $T2 -> T3$. There is a cycle, hence it is not serializable.

- (b) **(6 points)** Suppose you are using a modified protocol called **1.5PL** for concurrency management as follows: Transactions must first get shared locks before reading an item and write locks before writing items (the locks can be obtained any time before the read/write operations). A transaction can get any new locks while it is in the growing phase. However, once a transaction releases any locks, it enters the shrinking phase. In the shrinking phase, transactions are only allowed to get read locks.

Does **1.5PL** guarantee serializability, and why or why not? You must give a detailed explanation for your answer. If it does not, show an example of a schedule that is not serializable but possible under this protocol.

Solutions: It does not guarantee serializability.

$w_2(Y)r_1(Y)w_1(Z)r_2(Z)$: this schedule is not serializable but is possible under this schedule.

Question 3 (16 points). Suppose you are given the following statistics for the relation:

`DataPoints(id, event_id, seqno, starttime, endtime, segment_id)`

`Tuples(DataPoints)=5,000,000` `Values(DataPoints.event_id)=50,000`

`Pages(DataPoints)=50,000` `Values(DataPoints.segment_id)=100`

`Values(DataPoints.starttime)=400` between 05:00 and 22:00

Index I1 (height 2) on `DataPoints(segment_id, event_id)` with 5,000 leaf nodes

(a) What is the estimated cost of answering Q1 using I1?

Q1: `select starttime, endtime from datapoints where event_id = 123`

(b) What is the estimated cost of answering Q2 using I1?

Q2: `select starttime, endtime from datapoints where segment_id = 456 and event_id = 123`

(c) What is the estimated cost of answering Q3 using I1?

Q3: `select event_id from datapoints where segment_id = 456 and starttime > time '17:00'`

(d) Suppose Q1, Q2, Q3 are asked with equal frequency. Is I1 the best index to create to help with all these queries or should you create alternate one or more indices instead? Discuss why or why not.

Solutions:

(a) Q1: Selectivity: $1/50,000$, Tuples: $5,000,000/50,000 = 100$

scan all leaf nodes, plus two internal, 5002, read matching 100 tuples, 5,102 total.

(b) Q2: Selectivity: $1/(50,000*100)$, Tuples = 1

expect 1 tuple to match, scan 3 nodes (1 leaf) plus 1 for the table.

(c) Q3: Selectivity for the index = $1/100$ (expect $5,000,000/100 = 50,000$ tuples to match)

50K tuples to read but scan for $5000/100 = 50$ nodes + 2.

(d) It seems `event_id` is a lot more selective, so I would go for an index on `event_id` and another index on `segment_id`. I can also consider the second index on `segment_id, starttime` but `starttime` is not very selective.

Question 4 (12 points). For the following, show all your computations and explain your solutions with a short sentence.

- (a) What is the cost of block nested join with $M=101$ of R (left or outer join) and S (right or inner relation) given $PAGES(R) = 400$ and $PAGES(S) = 2000$?
- (b) What is the cost of sorting of a relation T with $M=101$ where $PAGES(T) = 50,000$?
- (c) Suppose now $T = R \bowtie S$ and we are pipelining the output of the join from part (a) to the sort in part (b). The output of part (b) is piped into a group by that is computed on the fly ($M=1$). What is the total cost of this query plan?

Solutions:

(a) Join: Read R once and S , $400/100=4$ times. Total cost: $400+2000*4=8,400$.

(b) Sort: Pass 1 100,000: produces 500 sorted groups, Pass 2 100,000: produces 5 sorted groups, Pass 2 50,000 completes the sort.

250,000 pages total.

(c) Combined: First read of T is not needed and no cost to group by, total: $8,400+ 200,000=208,400$.

Question 5 (10 points). Suppose you are given the following contents of the log and data pages on disk after a crash and there is no other log entry. Answer the following questions yes/no/not enough information with a short explanation as to why or why not.

LOG:

LSN	LOG ENTRY	PREVLSN
1	T1 update P1 10 20	-
2	T2 update P2 A B	-
3	T3 update P4 X Y	-
4	T3 update P3 AA BB	4
5	T1 update P2 B C	1
6	commit T3	
7	T2 update P5 55 66	2
8	T2 update P1 20 30	7
9	commit T1	5

DATA PAGES:

PAGE ID	LSN OF LAST UPDATE
P1	1
P2	5
P3	4
P4	3
P5	-

- Is STEAL used? Explain why or why not.
- Is FORCE used? Explain why or why not.
- Which operations need to be REDOne and UNDOne? List the LSNs and discuss why.

Solutions: (a) There is no instance of STEAL, but we cannot conclude no STEAL (not enough information). (We will also accept that there is STEAL because we have actually written to disk changes by LSN=2 to P2 which was overwritten by T1.)

(b) All pages modified by committed transaction are written to disk, so it is possible FORCE is used but we cannot conclude for sure.

(c) No REDO or UNDO necessary. (This part is not needed but there is a concern with log entry 2 which was later overwritten by a committed transaction (for P2). Technically we must UNDO it but we don't want to UNDO the changes by an committed transaction. Basically this operation should not have been allowed.)

Question 6 (12 points). For each new relation given below, answer the following:

- (a) Write down the associated functional dependencies corresponding to the described rules.
- (b) Is this relation in 3NF? Is it in BCNF? Show why or why not.
- (c) If the relation is not in 3NF, then compute a 3NF decomposition for it.

UserGroups(groupname, URL, owner, membername, memberrole, startdate)

A user group (groupname) has a single owner and URL, and multiple members (membername). Each member may have multiple roles in the group, but for each role for a specific group, there is a single member and a single start date.

Solutions:

groupname -> owner, URL
groupname, role -> membername, startdate

key, groupname, role. Not in 3NF or BCNF.

(groupname, owner, url)
(groupname, role, membername, startdate)

UserNotes(username, event_id, note_text, note_id)

Each user can take many notes for an event. For each note_id, there is a unique username, event_id and note_text.

Solutions:

note_id -> username, event_id, note_text

Key note_id

Both in 3NF and BCNF.

Profile(username, popular_segment, popular_starttime, popular_landmark, best_friend, eventtype)

For each username, there is a single popular_starttime, but multiple popular_segment and popular_landmarks. For each popular_segment, username and best_friend, there is an eventtype.

Solutions:

username -> popular_starttime
username, popular_segment, best_friend -> event_type

Key: username, popular_segment, best_friend, popular_landmark

Not in BCNF, 3NF.

username, popular_starttime
username, popular_segment, best_friend, event_type
username, popular_segment, best_friend, popular_landmark

This page is left blank for scratch work, random thoughts and pictures!

Congratulations, DBS is done.

Let me know if you wish to be a mentor for DBS next semester. It is a great learning experience.

Also, have a great summer.



Data model to be used in Final Exam

```
create table users ( -- all users in the system
    username      varchar(12) primary key
    , name        varchar(100) not null
    , email       varchar(100) not null
    , password    varchar(100) not null
    , address     varchar(100)
) ;

create table friendswith (
    -- friendship is mutual, but stored in one direction only,
    -- username1 is the person initiated the friendship
    username1     varchar(12)
    , username2   varchar(12)
    , sincewhen   date -- when friendship was confirmed
    , primary key (username1, username2)
    , foreign key (username1) references users(username)
    , foreign key (username2) references users(username)
) ;

create table landmarks (
    id            int primary key
    , landmarkname varchar(100) not null
    , landmarktype varchar(100) --e.g. building, monument, etc.
    , state       varchar(100)
    , city        varchar(100)
    , zip         varchar(20)
    , country     varchar(100)
) ;

create table segments (
    id            int primary key
    , startx      numeric(8,4) not null
    , starty      numeric(8,4) not null
    , endx        numeric(8,4) not null
    , endy        numeric(8,4) not null
    , landmark_id int
    , foreign key (landmark_id) references landmarks(id)
) ;

create table events (
    id            int primary key
    , username     varchar(12) not null
    , eventtype    varchar(100) not null --cycling, running, etc.
    , eventdate    date not null
    , starttime    time not null
    , endtime      time
    , foreign key (username) references users(username)
) ;

create table datapoints (
    id            int primary key
    , event_id     int not null
```

```

    , seqno            int not null
    , starttime        time
    , endtime          time
    , segment_id       int
    , foreign key (event_id) references events(id)
    , foreign key (segment_id) references segments(id)
) ;
-- a data point is a segment on someone's running or cycling event,
-- the first data point has seqno=1, followed by seqno: 2,3,4 ...
-- describing the segments that one has passed during the event.
-- The end point of a segment for a data point with seqno x is the starting
-- point of the segment for the data point with seqno x+1.

create table comments (
    -- each comment is made by the user with given username
    -- for a specific event
    id                int primary key
    , username        varchar(12) not null
    , event_id        int not null
    , comment_text    text not null
    , whenposted      date
    , foreign key (username) references users(username)
    , foreign key (event_id) references events(id)
) ;

```